

Chapter 13: I/O Systems



Chapter 13: I/O Systems

- ❑ I/O Hardware
- ❑ Application I/O Interface
- ❑ Kernel I/O Subsystem
- ❑ Transforming I/O Requests to Hardware Operations
- ❑ Streams
- ❑ Performance



Objectives

- Explore the structure of an operating system's I/O subsystem
- Discuss the principles of I/O hardware and its complexity
- Provide details of the performance aspects of I/O hardware and software

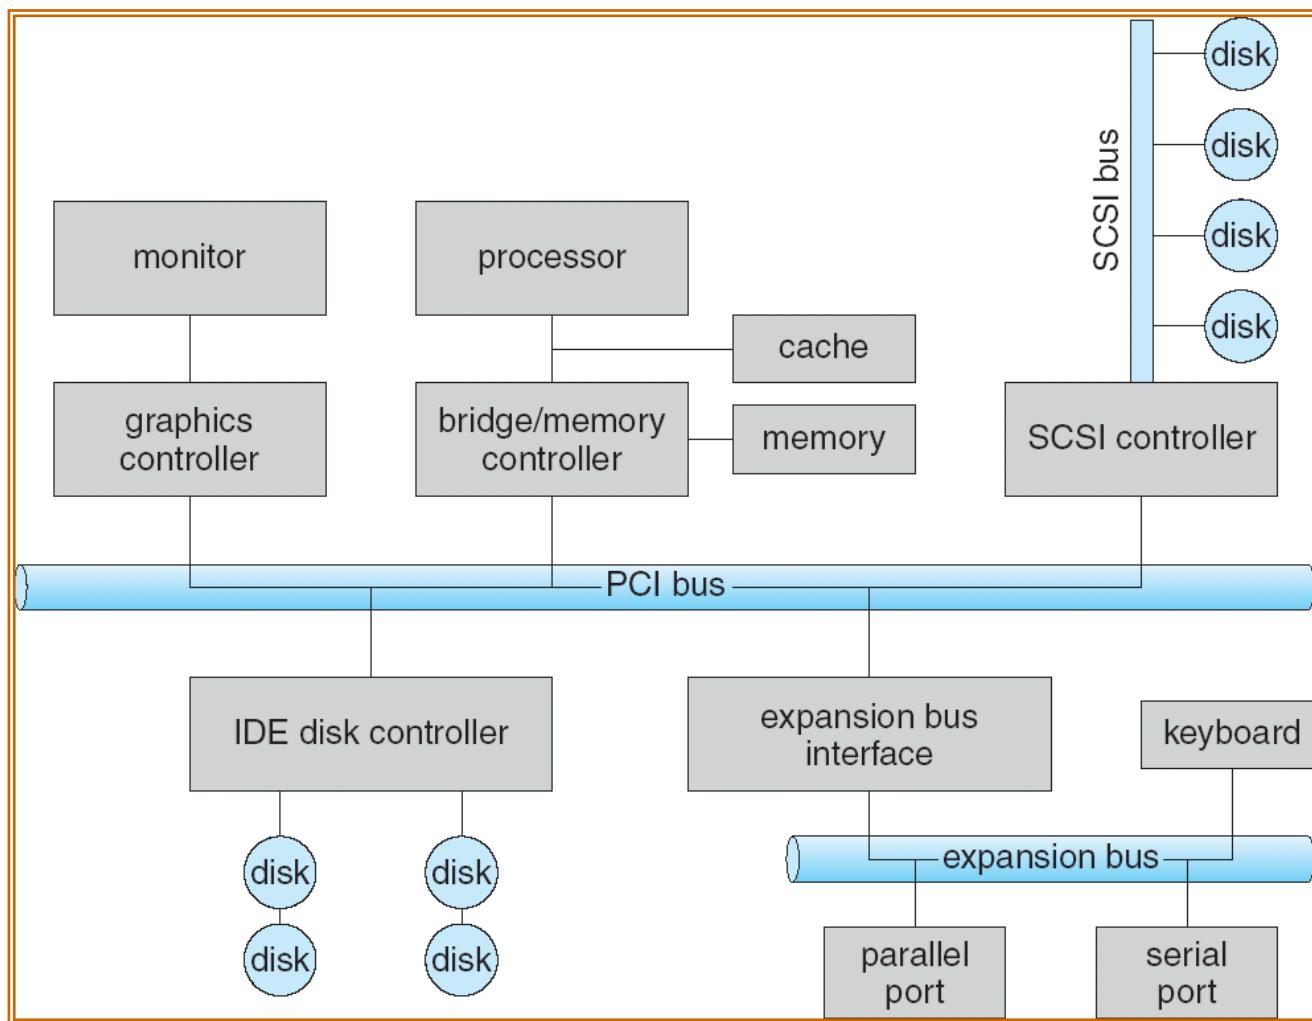


I/O Hardware

- Incredible variety of I/O devices
- Common concepts
 - Port
 - Bus (daisy chain or shared direct access)
 - Controller (host adapter)
- I/O instructions control devices
- Devices have addresses, used by
 - Direct I/O instructions (专有I/O通道)
 - Memory-mapped I/O (将I/O映射到内存里面,从而使得I/O和内存管理得到统一)

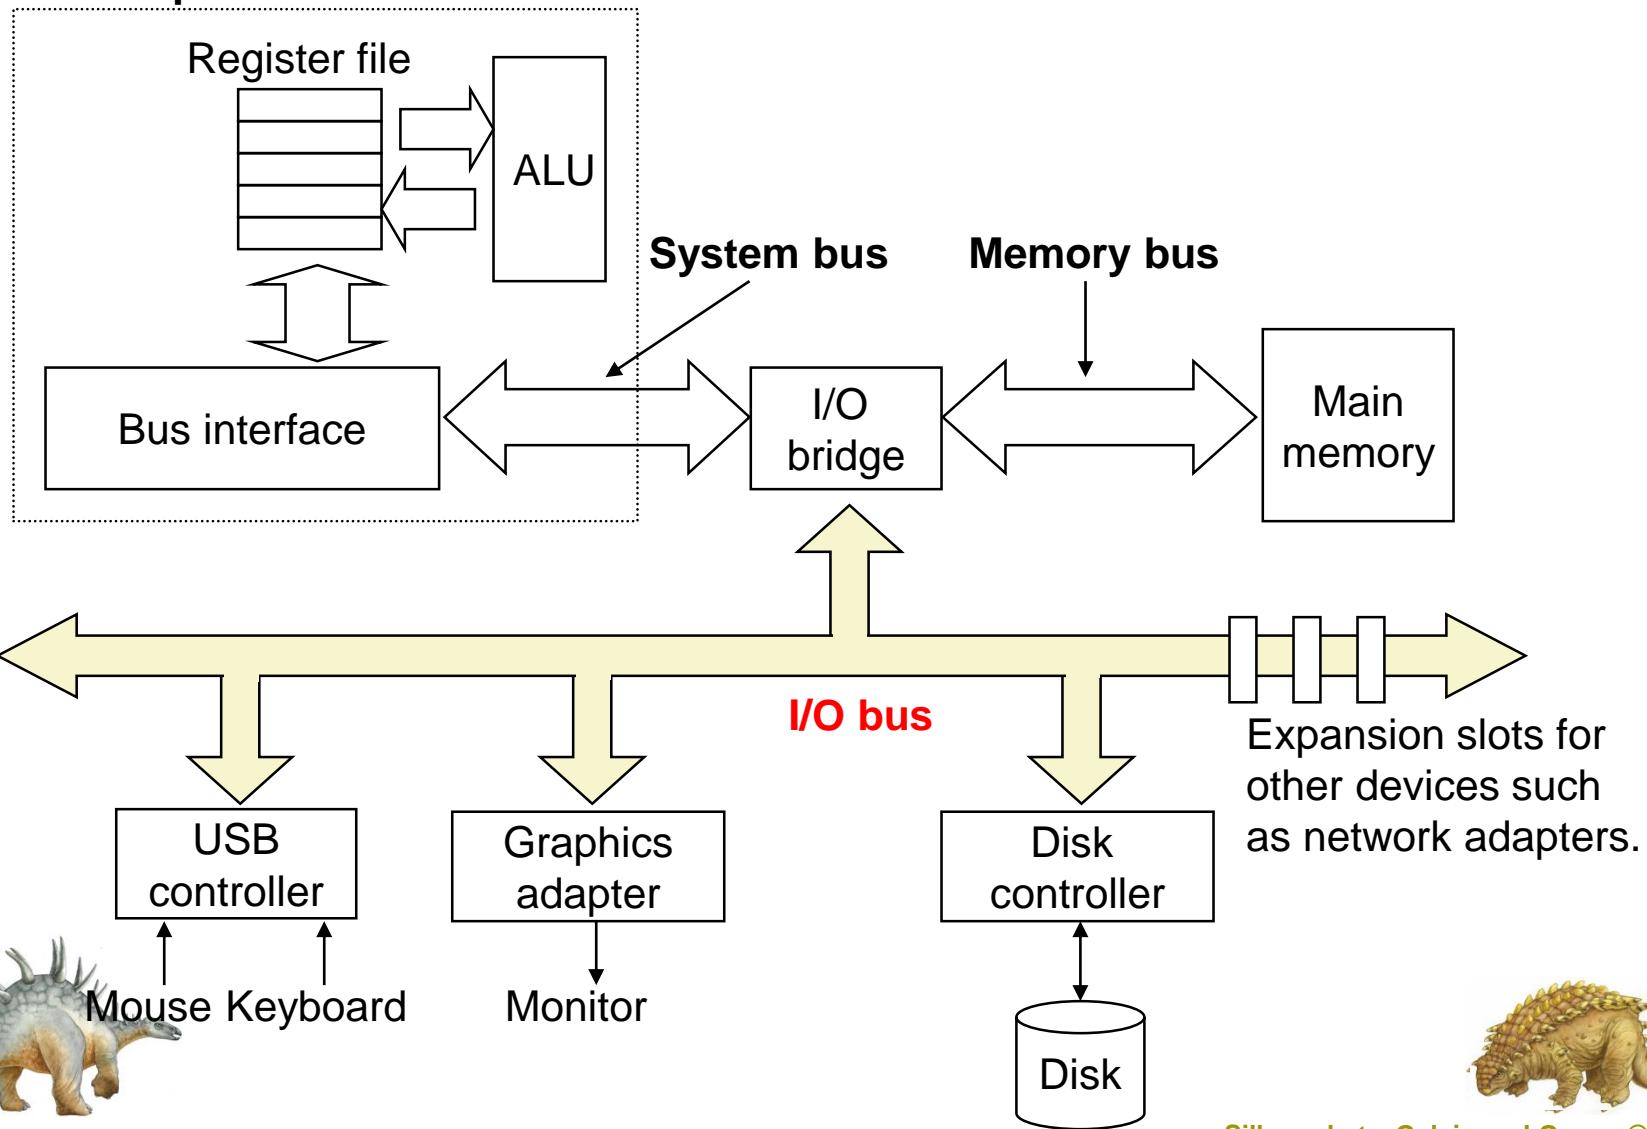


A Typical PC Bus Structure



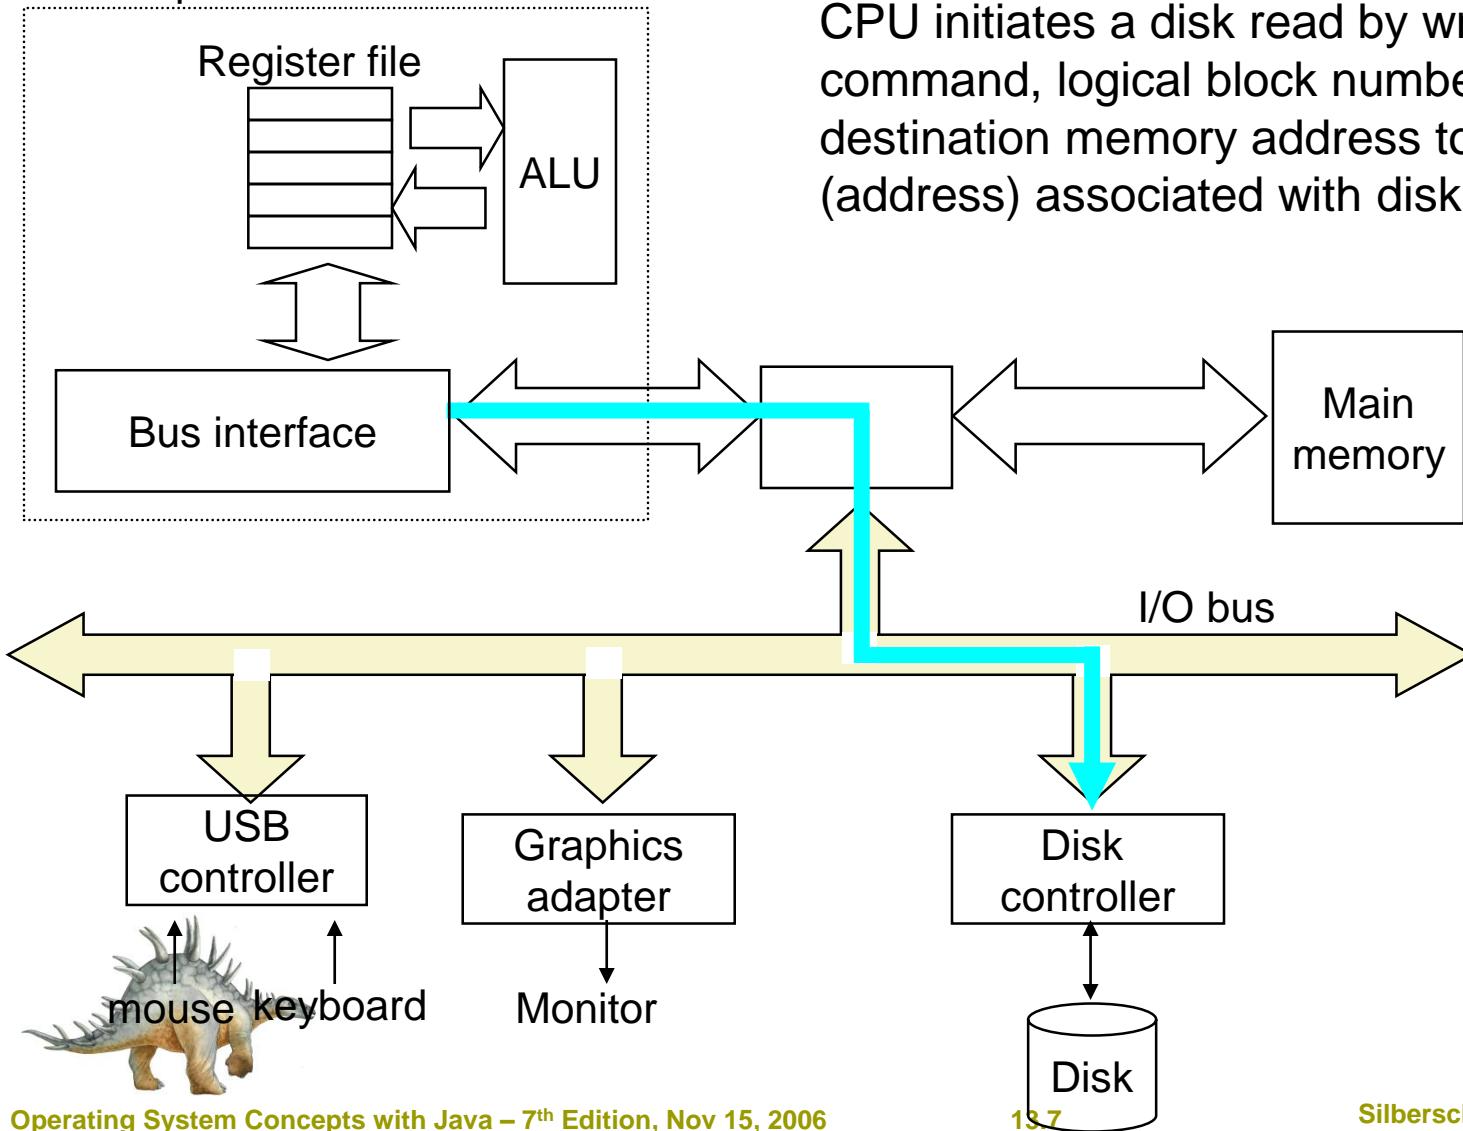
I/O Bus

CPU chip



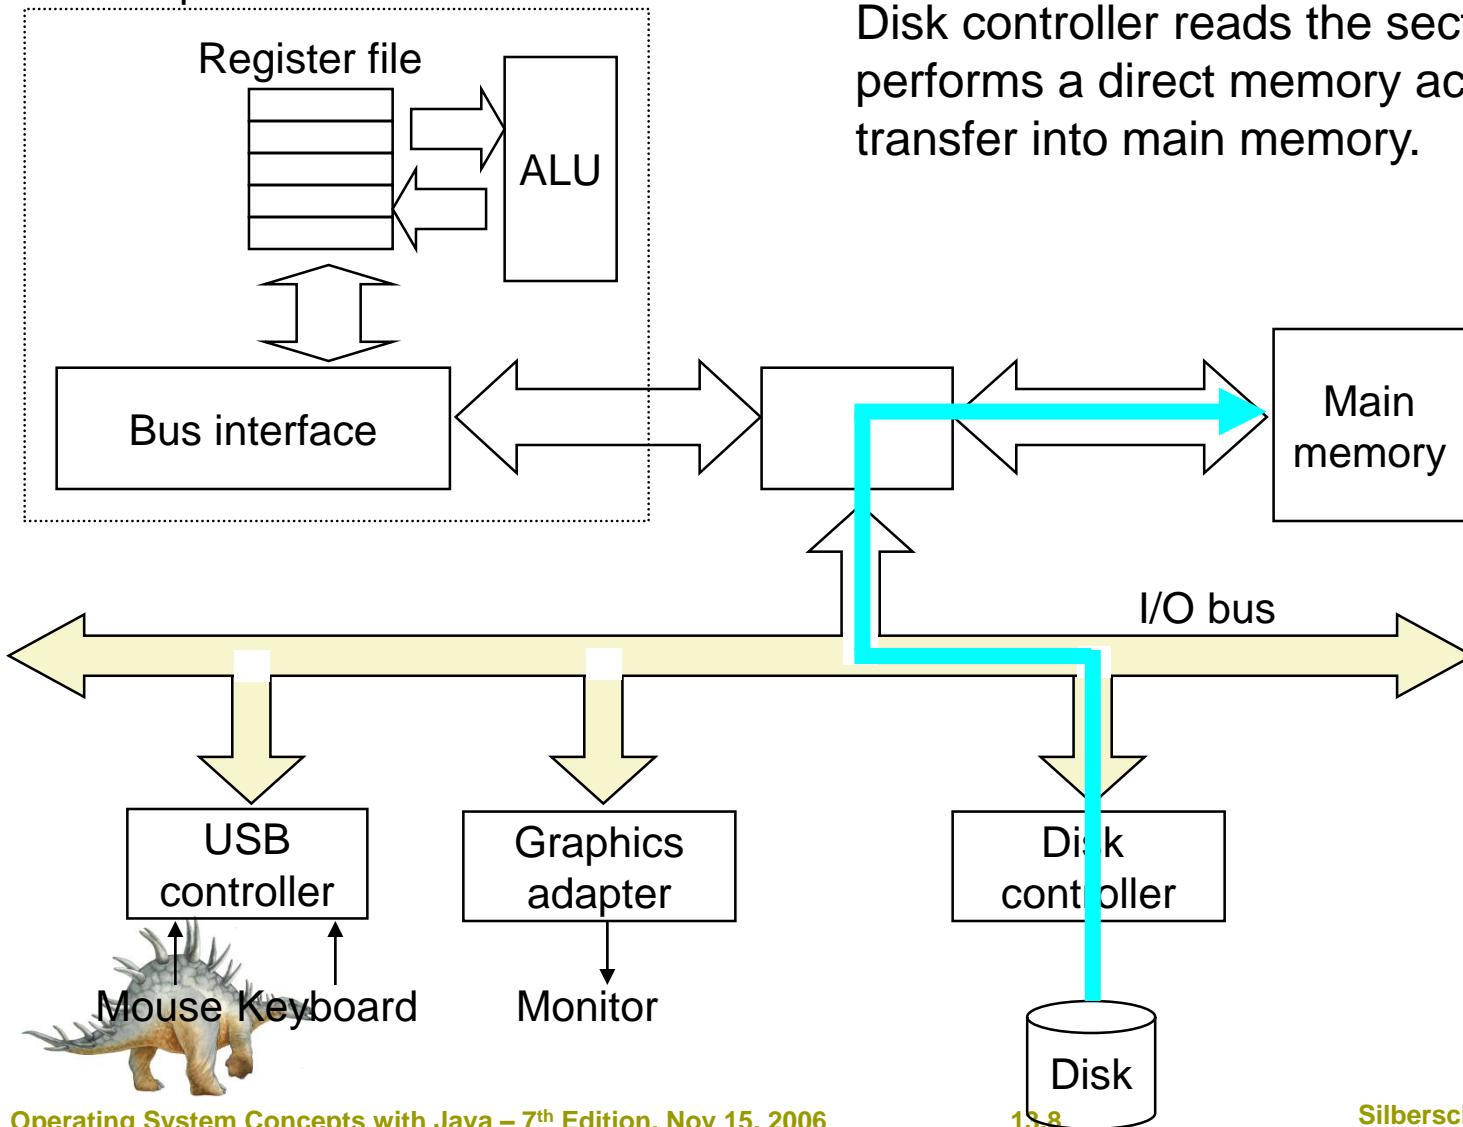
Reading a Disk Sector (1)

CPU chip



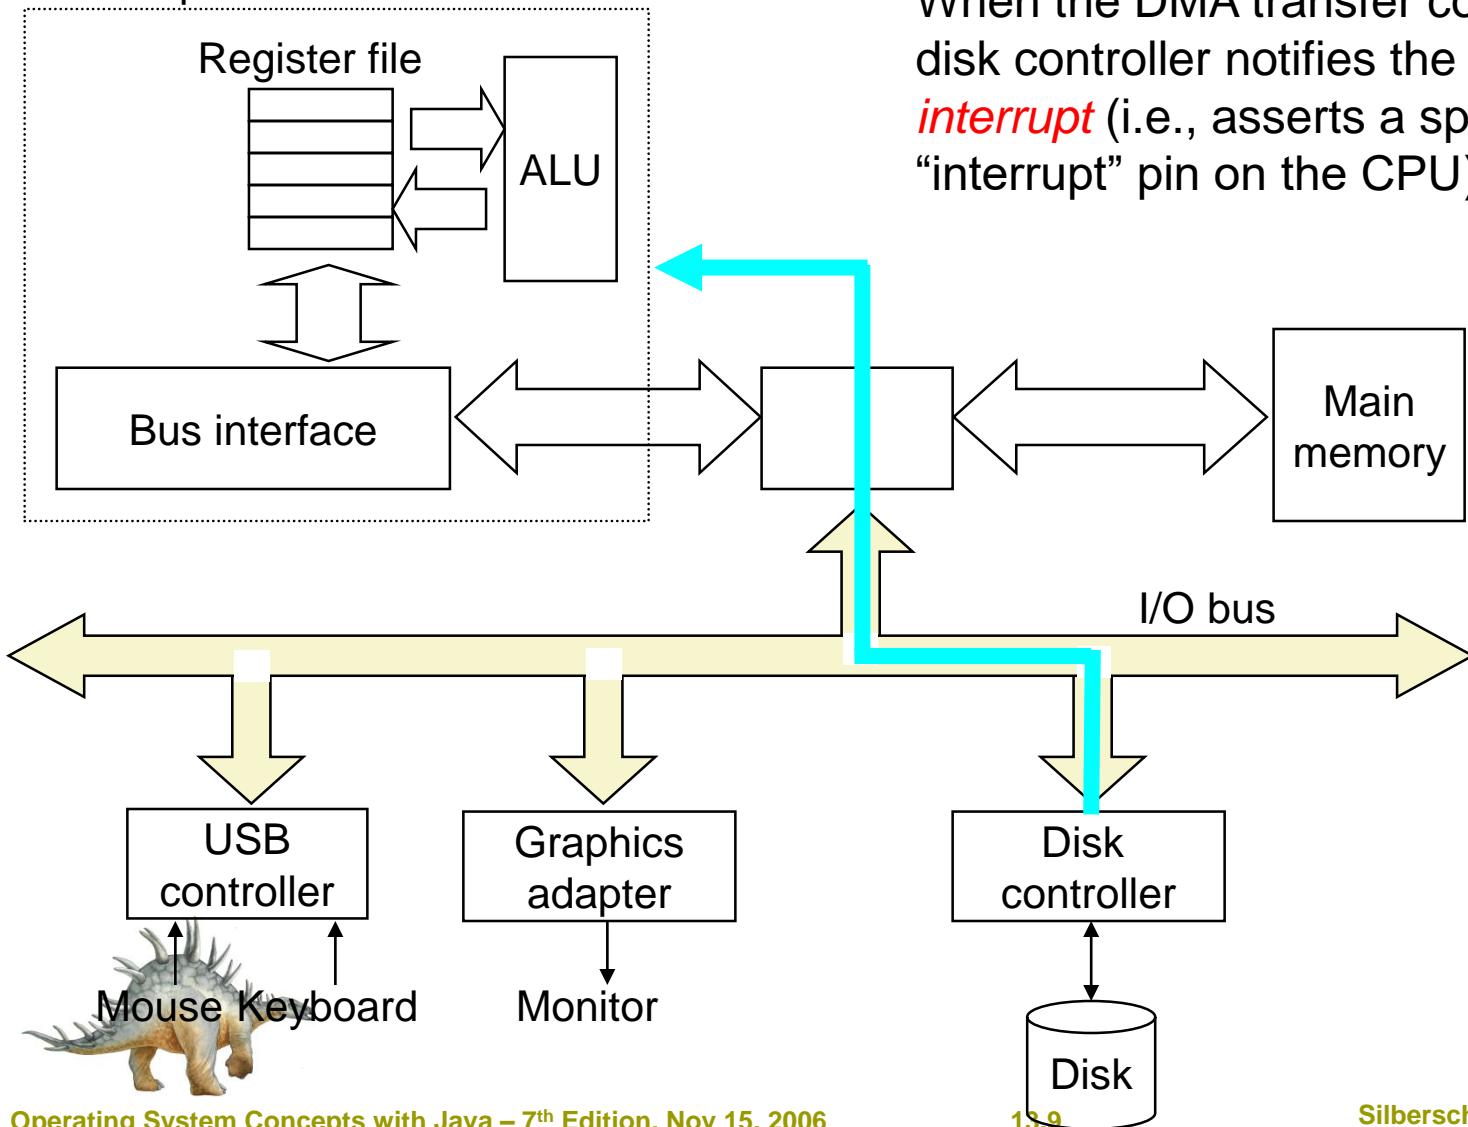
Reading a Disk Sector (2)

CPU chip



Reading a Disk Sector (3)

CPU chip



When the DMA transfer completes, the disk controller notifies the CPU with an *interrupt* (i.e., asserts a special “interrupt” pin on the CPU)

Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)



Polling

- Determines **state of device**
 - command-ready
 - busy
 - Error
- **Busy-wait** cycle to wait for I/O from device

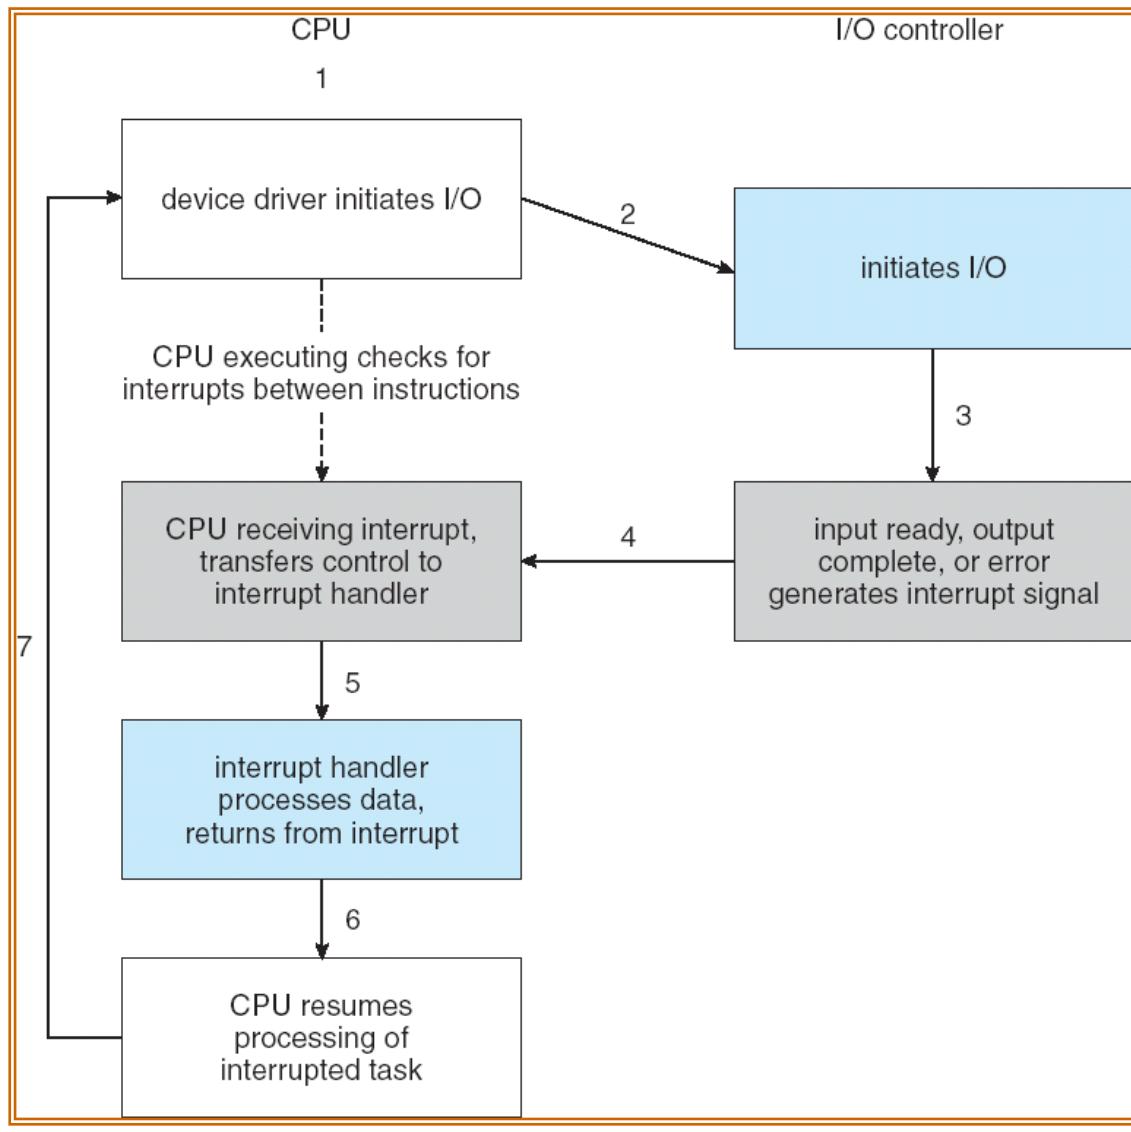


Interrupts

- ❑ CPU **Interrupt-request line** triggered by I/O device
- ❑ **Interrupt handler** receives interrupts
- ❑ **Maskable** to ignore or delay some interrupts
- ❑ Interrupt vector to dispatch interrupt to correct handler
 - Based on priority
 - Some **nonmaskable**
- ❑ Interrupt mechanism also used for exceptions



Interrupt-Driven I/O Cycle



Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts

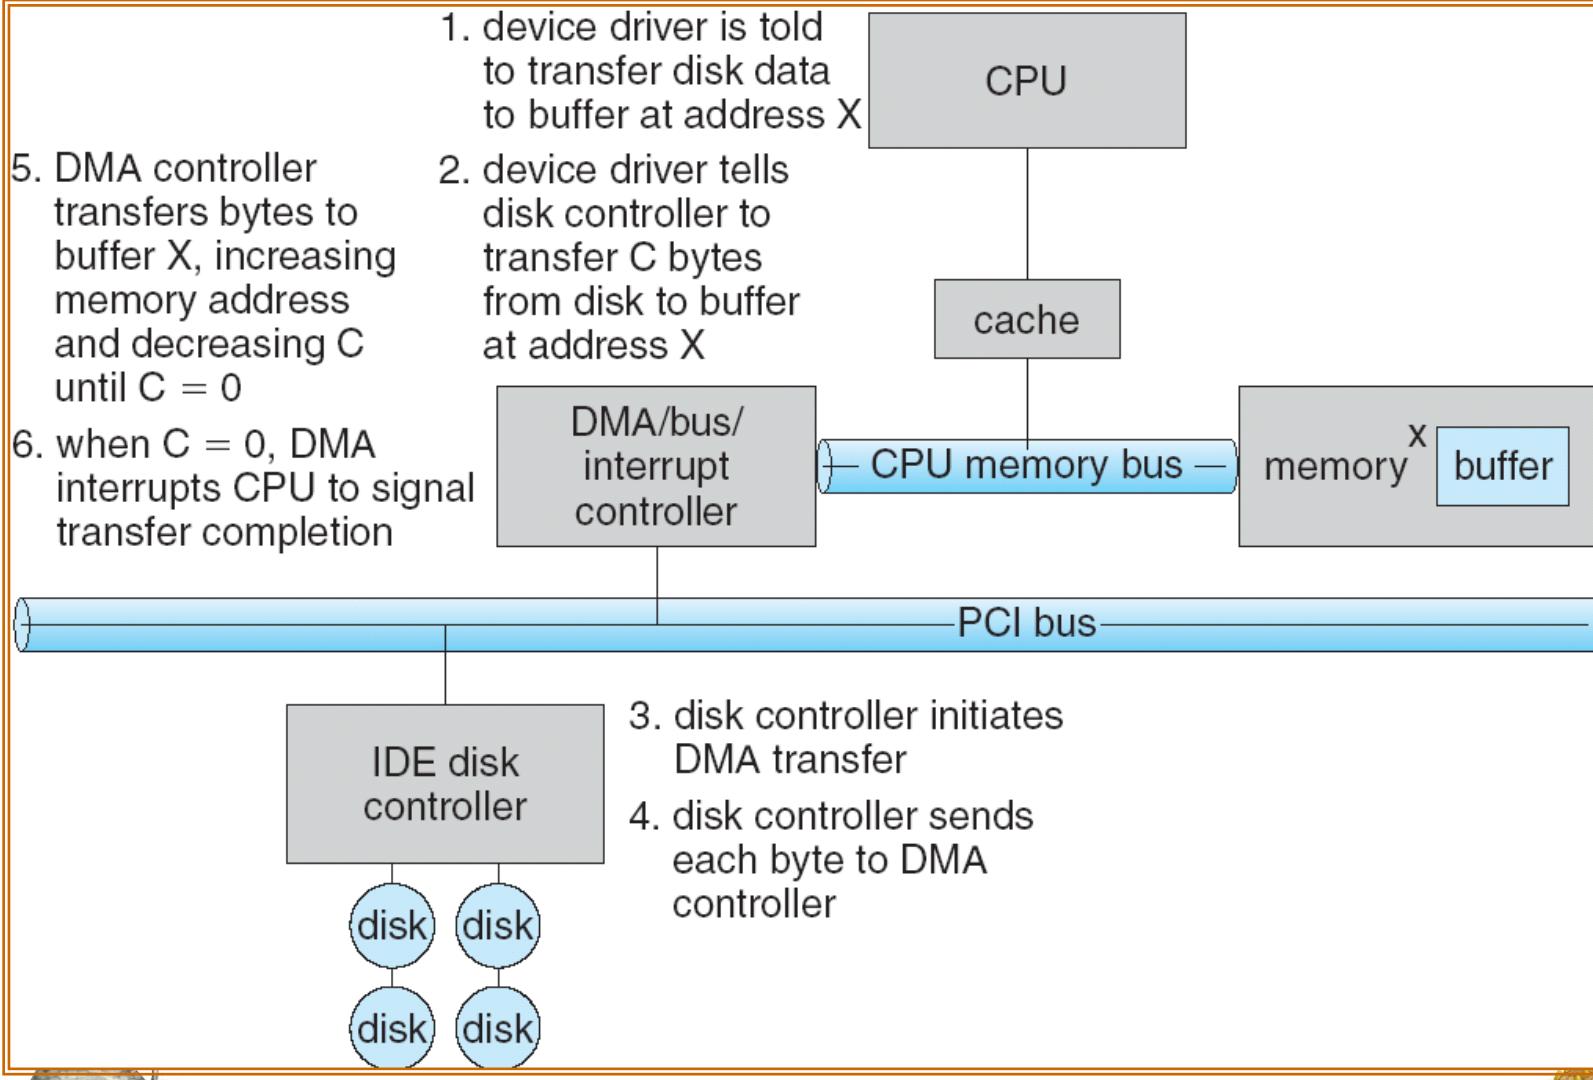


Direct Memory Access

- Used to avoid **programmed I/O** for large data movement
- Requires **DMA** controller
- Bypasses CPU to transfer data directly between I/O device and memory



Six Step Process to Perform DMA Transfer

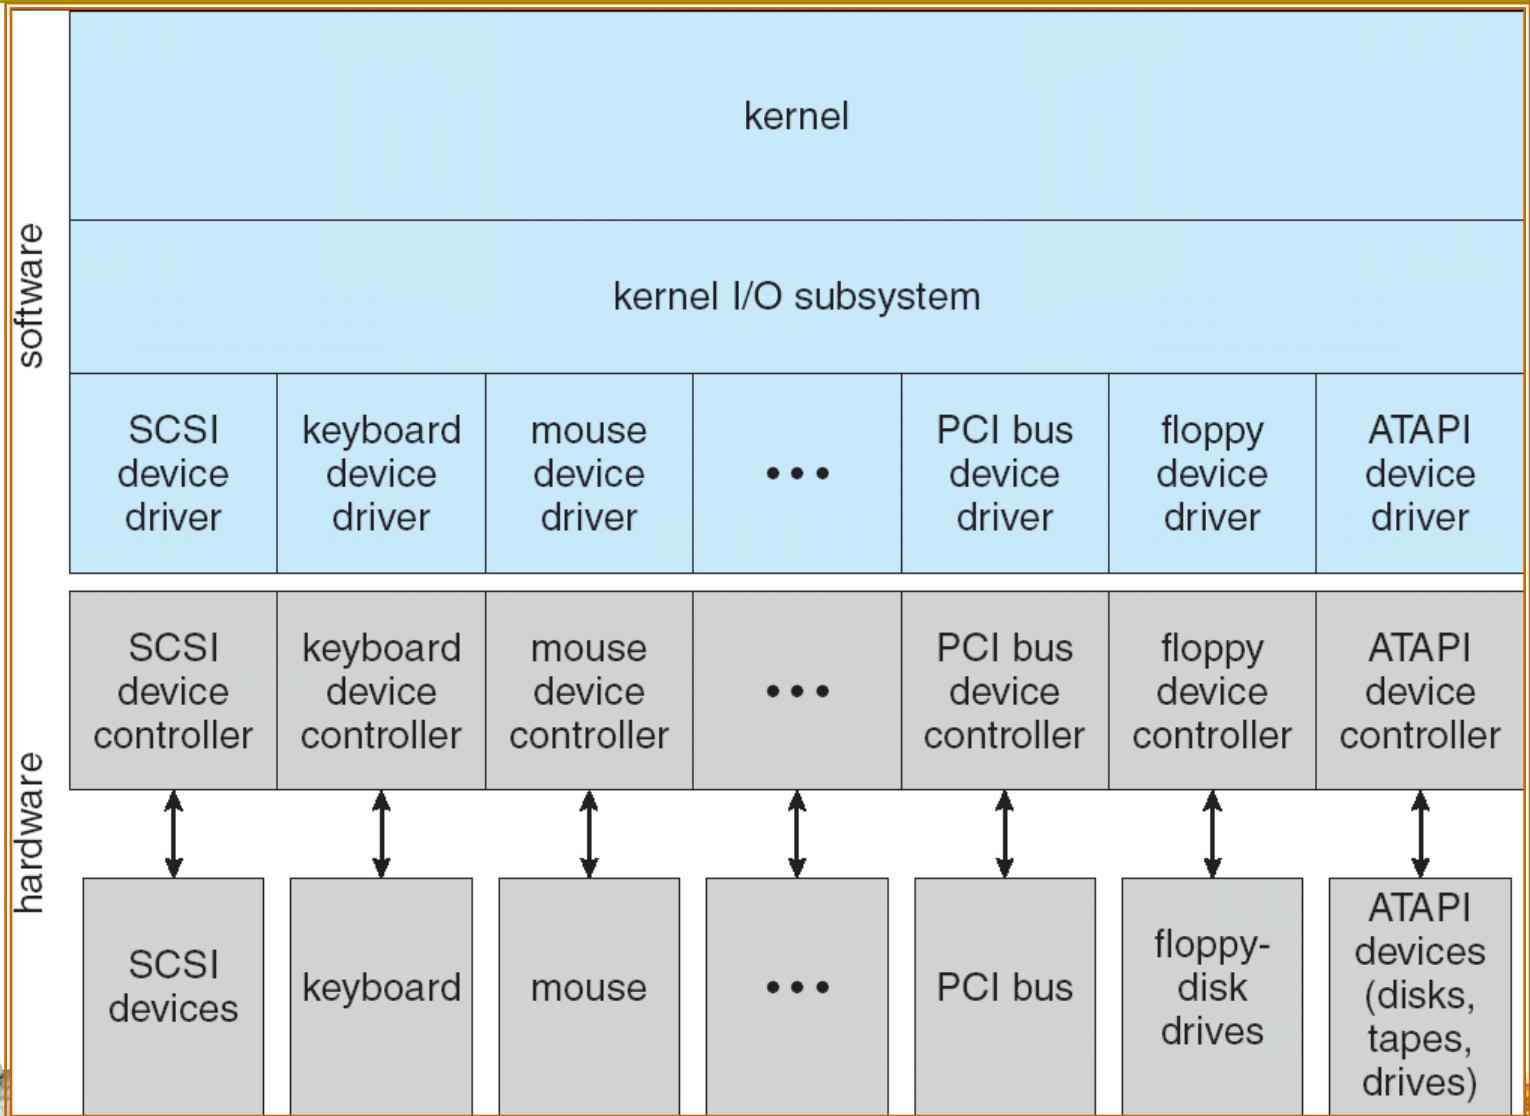


Application I/O Interface

- I/O system calls **encapsulate** device behaviors in generic classes
- Device-driver layer **hides differences** among I/O controllers from kernel
- Devices vary in many dimensions
 - **Character-stream or block**
 - **Sequential or random-access**
 - **Sharable or dedicated**
 - **Speed of operation**
 - **read-write, read only, or write only**



A Kernel I/O Structure



Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read-write	CD-ROM graphics controller disk



Block and Character Devices

- **Block devices** include disk drives
 - Commands include **read**, **write**, **seek**
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- **Character devices** include keyboards, mice, serial ports
 - Commands include **get**, **put**
 - Libraries layered on top allow line editing



Network Devices

- ❑ Varying enough from block and character to have own interface
- ❑ Unix and Windows NT/9x/2000 include socket interface
 - Separates network protocol from network operation
 - Includes select functionality
- ❑ Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)



Clocks and Timers

- Provide current time, elapsed time, timer
- **Programmable interval timer** used for timings, periodic interrupts
- ioctl (on UNIX) covers odd aspects of I/O such as clocks and timers

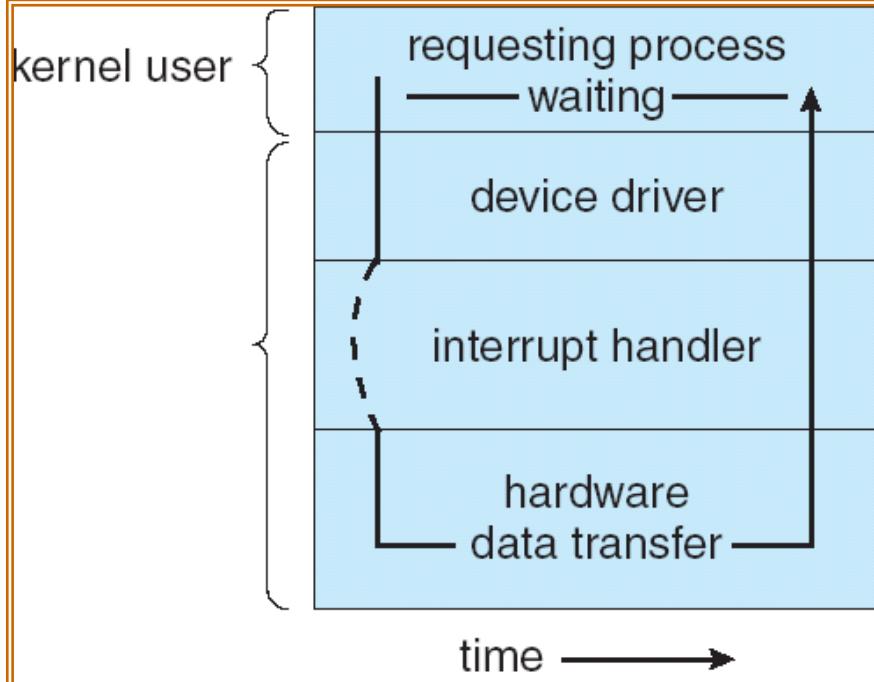


Blocking and Nonblocking I/O

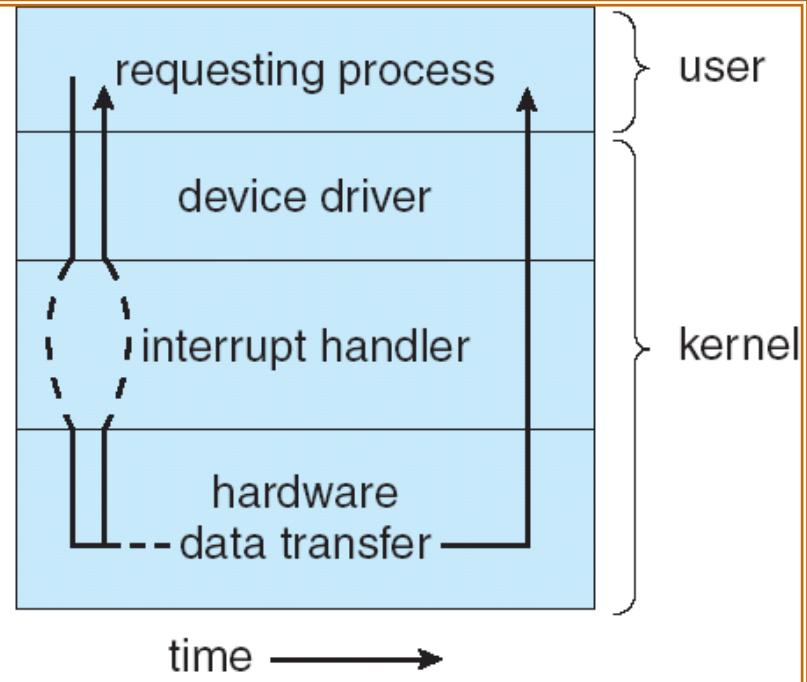
- **Blocking** - process **suspended** until I/O completed
 - Easy to use and understand
 - Insufficient for some needs
- **Nonblocking** - I/O call returns as much as available
 - User interface, data copy (buffered I/O)
 - Implemented via multi-threading
 - Returns quickly with **count of bytes read or written**
- **Asynchronous** - process runs while I/O executes
 - Difficult to use
 - I/O subsystem signals process when I/O completed



Two I/O Methods



Synchronous



Asynchronous

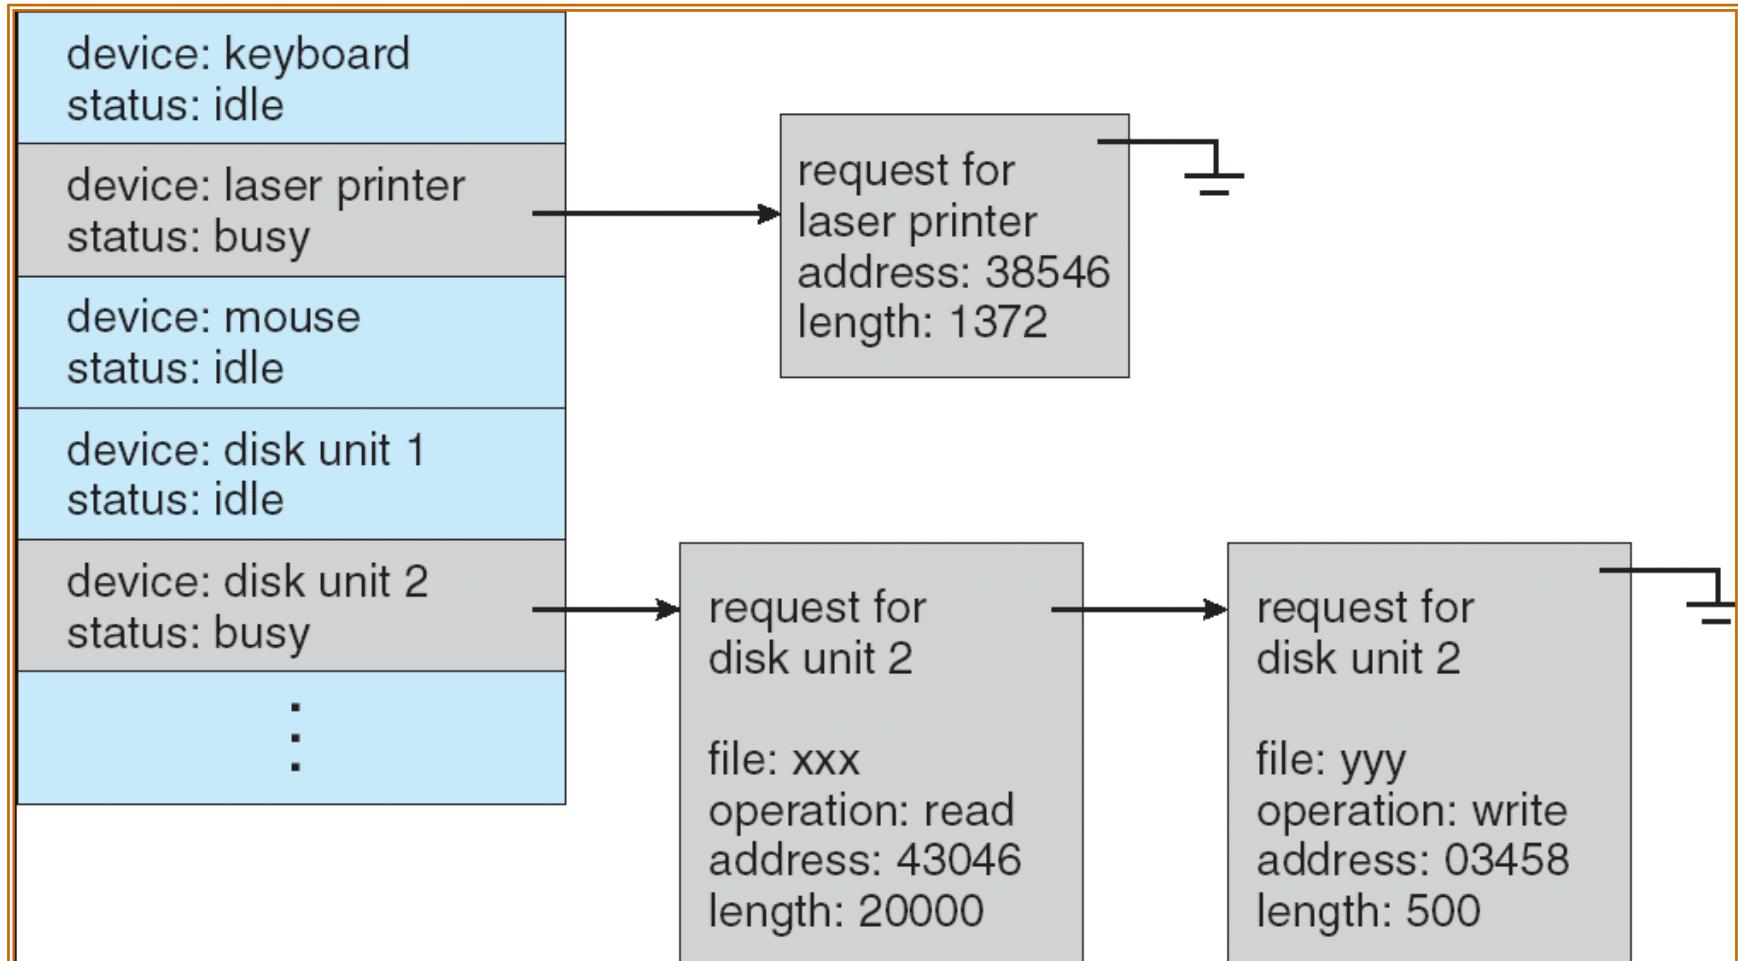


Kernel I/O Subsystem

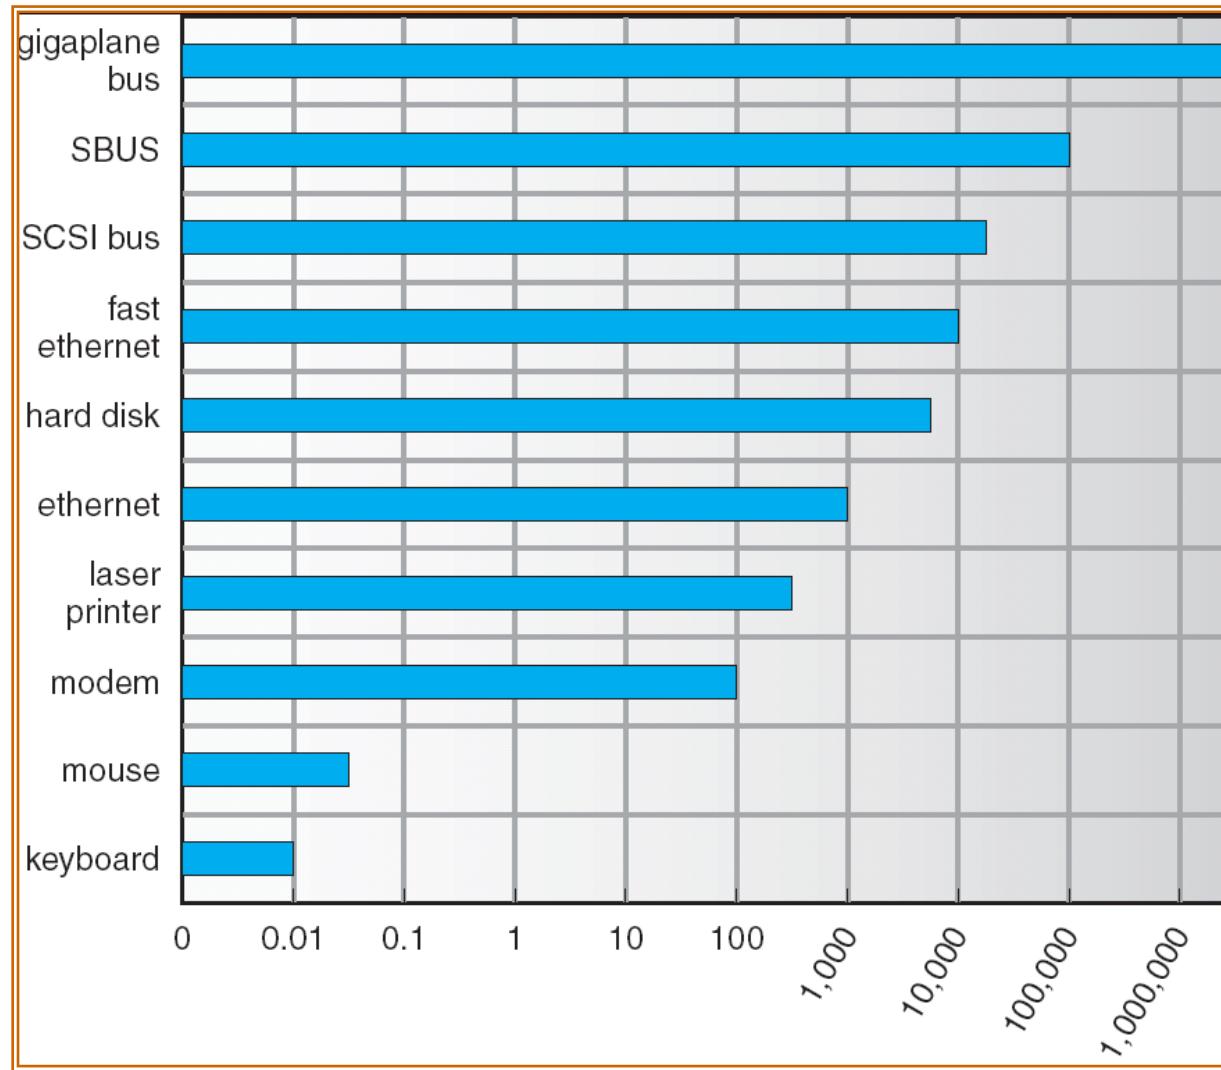
- **Scheduling**
 - Some I/O request ordering via per-device queue
 - Some OSs try fairness
- **Buffering** - store data **in memory** while transferring between devices
 - To cope with device speed **mismatch**
 - To cope with device transfer **size mismatch**
 - **To maintain “copy semantics”**



Device-status Table



Sun Enterprise 6000 Device-Transfer Rates



Kernel I/O Subsystem

- **Caching** - fast memory holding copy of data
 - Always just a copy (buffer)
 - Key to performance
- **Spooling** - hold output for a device
 - If device can serve only one request at a time
 - i.e., Printing
- **Device reservation** - provides exclusive access to a device
 - System calls for allocation and deallocation
 - Watch out for deadlock



Error Handling

- OS can recover from disk read, device unavailable, transient write failures
- Most return an **error number** or code when I/O request fails
- System **error logs** hold problem reports



Improving Performance

- Reduce number of context switches
- Reduce data copying
- Reduce interrupts by using large transfers,
smart controllers, polling
- Use DMA
- Balance CPU, memory, bus, and I/O
performance for highest throughput



End of Chapter 13



I/O管理

设备管理是指操作系统对除了CPU和内存以外的所有输入/输出设备的管理，诸如设备控制器、通道、中断控制器等等，因此为了提高计算机系统的整体效率，除了需要对中央处理器CPU合理调度、对内存合理使用之外，对系统中的设备也要实施数行之有效的管理，这样才能真正发挥计算机系统的整体效率。



I/O外设分类

人可读；

- 计算机用户间的交互；
- 打印机；
- 视频终端；
- 键盘、鼠标；

机器可读；

- 电子设备间通讯；
- 磁盘、磁带设备；
- 传感器；
- 控制器；
- 传动器；



通信；

- 远程设备间通信；
- 数字线路驱动器；
- 调制解调器；



I/O功能的组织

➤ 可编程I/O

处理器代表进程给I/O模块发送一个I/O命令，进程进入忙等待，操作完成后，才可继续执行；

➤ 中断驱动I/O

➤ 直接存储访问（DMA）

DMA模块控制主存和I/O模块间数据交换，只有当整个数据块结束后，处理器才被中断；



I/O功能的发展

1. 处理器直接控制外部设备：

由 CPU 负责在机器内存与设备控制器数据寄存器之间进行数据
传送

2. 非中断可编程I/O

3. 中断驱动I / O

为了减少设备驱动程序不断地询问控制器状态寄存器的开销当
I/O 操作结束后，由设备控制器主动通知设备驱动程序

4. DMA

5. 通道

6. I/O处理器



中断方式传送

为了减少CPU测试等待时间和CPU与外设的并行工作能力，引入了中断控制传输方式。

- 进程CPU发START命令启动外设准备数据；
- 进程放弃CPU，进程调度选择其它进程占有CPU；
- 外设准备好数据，通过**中断控制器**向CPU发中断请求信号，CPU接到该信号且**在允许中断的情况下**，**响应该中断请求**，转**中断处理程序**，在**中断处理程序中完成数据输入。**



DMA方式

存储器直接存取方式；

内存和外设之间开辟直接的数据交换通路，由DMA控制器完成数据交换。

DMA方式 在传送开始需要CPU做一些初始化和传输结束做一些善后处理工作之外，在整个数据传输过程中，不需要CPU任何干预。

1. 进程要求输入数据，CPU将输入数据的内存地址及个数送入相应寄存器；
2. 进程进入等待状态，进程调度程序调度其他进程；
3. 在DMA控制器控制下，将数据寄存器的数据不断写入内存；



DMA方式

4. 传送完毕，DMA控制器向CPU发出中断请求，CPU转中断处理程序；
5. 中断处理程序结束，CPU返回被中断的进程继续执行。



DMA方式

1. 负责在高速外围设备与内存之间成批量的数据传输工作，但是不对数据作再加工处理，I/O操作类型简单；
2. 需要使用专门的DMA控制器：控制、状态寄存器、传送字数寄存器、内存地址寄存器和数据缓冲寄存器。
3. 采用“**偷窃总线控制权**”，不用CPU干预
4. 每传送一个数据并不产生中断，只有本次DMA传送的数据全部完毕时，才产生中断。



DMA方式与中断的主要区别

- 中断方式是在数据缓冲寄存区满后，发中断请求，CPU 进行中断处理，DMA方式则是在所要求传送的数据块全部传送结束时要求 CPU 进行中断处理，大大减少了 CPU进行中断处理的次数。
- 中断方式的数据传送是由 CPU 控制完成的而 DMA方式则是在 DMA控制器的控制下不经过 CPU 控制完成的。



操作系统设计问题-设计目标

1. 提高使用效率:

提高CPU与外设、外设与外设之间的并行工作的能力。
。中断和通道技术的引入提供了这种并行性。

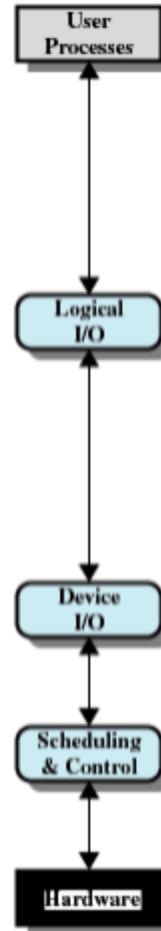
2. 通用性:

是使用户能独立于具体设备的复杂物理特性而方便的使用设备。也就是说避开各种物理设备的具体使用细节，呈现在用户面前的是简单、易操作的逻辑设备

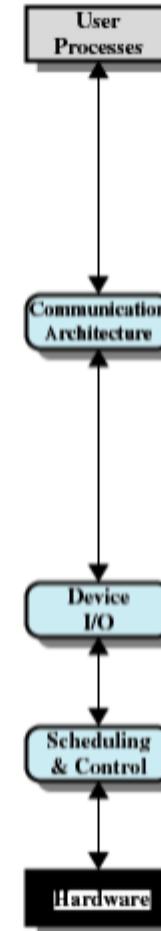


I/O功能逻辑结构

- **逻辑I/O:** 把设备当作一个逻辑资源来处理，并不关心实际控制设备的细节。
- **设备I/O:** 请求的操作和数据被转换成适当的I/O指令序列、通道命令和控制命令。
- **调度和控制:** 关于I/O操作的排队、调度和控制实际发生在这一层，该层与I/O模块和设备硬件真正发生交互的软件层。



(a) Local peripheral device



(b) Communications port



(c) File system

I/O缓冲

为解决计算机外设的和中央处理器之间速度不匹配，所产生的“瓶颈”现象，提高系统性能，因此引入了缓冲技术。

- 改善中央处理器与外围设备之间速度不配的矛盾，
- 提高CPU和I/O设备的并行性。

- 面向块的I/O设备；
- 面向流的I/O设备；



I/O缓冲-缓冲区管理

单缓冲

- 当用户进程发出I/O请求时，操作系统在主存的系统空间为该操作分配一个缓冲区，可以实现预读和滞后写。

双缓冲

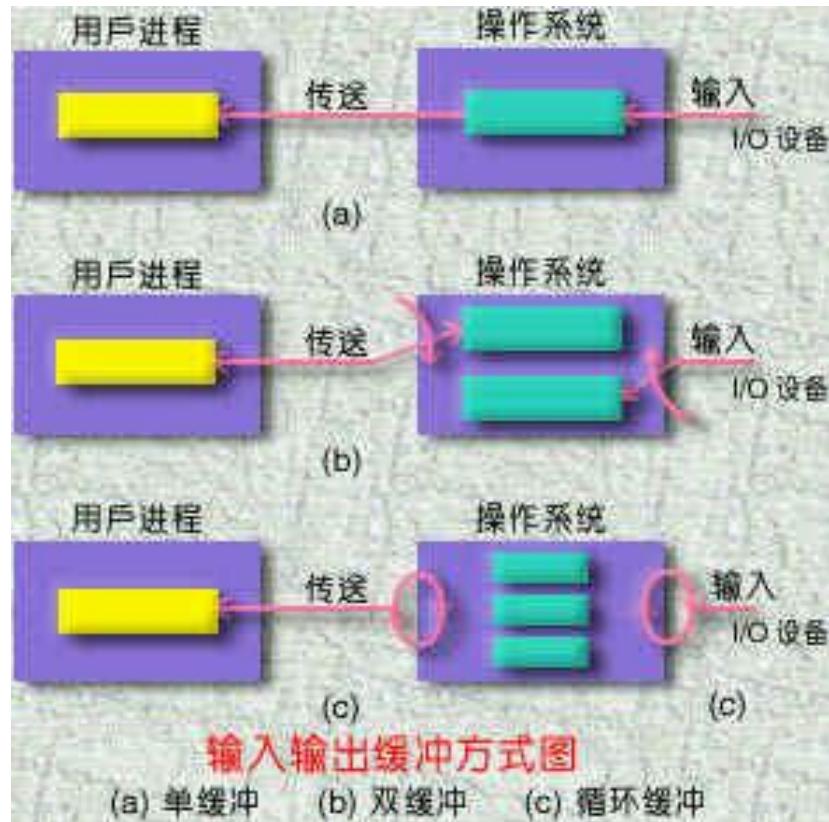
- 可以实现用户数据区—缓冲区之间交换数据和缓冲区—外设之间交换数据的并行。

循环缓冲

- 多个缓冲区连接起来统一管理



I/O缓冲-缓冲区管理



生产者/消费者模型
面向块：时间为 $\text{MAX}[C, T] + M$

生产者/消费者模型
面向块：时间为 $\text{MAX}[C, T]$

有界缓冲区
生产者/消费者模型

